

Rethinking how the industry approaches 'Chaos Engineering'

Info @ (Nora Jones)

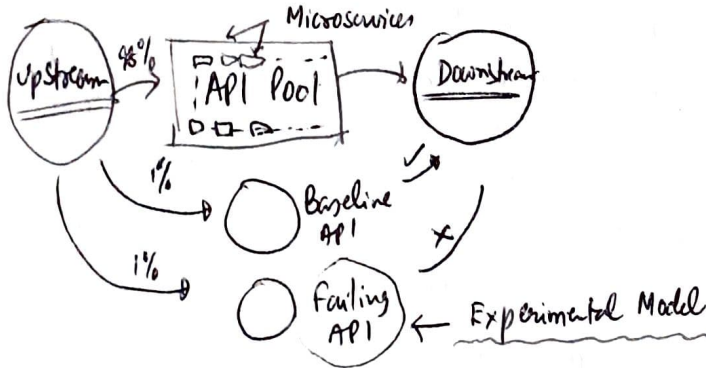
- Building resilience to withstand turbulent conditions in production
- Creating foresights & predictive maintenance for systems in production.
- Builds confidence in system behaviors through experiments.

Goals :


- Build toolings to stop issues / discover vulnerabilities in existing architecture.
- Build a culture of resilience to the unexpected.

ChAP → Automating production Chaos Experiment

[Chaos Automation in Production]



Building ChAP



Designing Meaningful Experiments : Who, What*, Why, When

- o "Apollo-1 Launch Rehearsal Test" (Novice Expert)
- o Digging advanced patterns in the systems / incident archive to build critical experimental modeling.

While building the exercise for experiments, a few important questions include →

- Which systems haven't failed in a while?
- Which system failure took us by surprise?
- Which incidents involved "unknown systems" or those that needed "experts" to step in & resolve?
- Which incidents involved people that haven't worked together much?
- What did recent "near misses" look like?
- Which incidents involved difficulties in figuring out what was even going on?

SCOPE

Steady-State (baseline) → Measurable output indicating normal behavior
 → Where we are injecting failure → diff. b/w "normal" & "good" operation.

Safe-guarding Experiments

Creating chaos is easy, what's difficult is :-
 Minimizing blast radius, safety-net implementation, & putting active listeners on when things go wrong.

(Post Experiment) Asking Better Experiments

- What / Why did you experiment?
- What was suspecting or new?
- What mental models got recalibrated?

Chaos Engineering can be used as cultural artifacts.

↳ Good way to learn about collaboration & system characteristics.

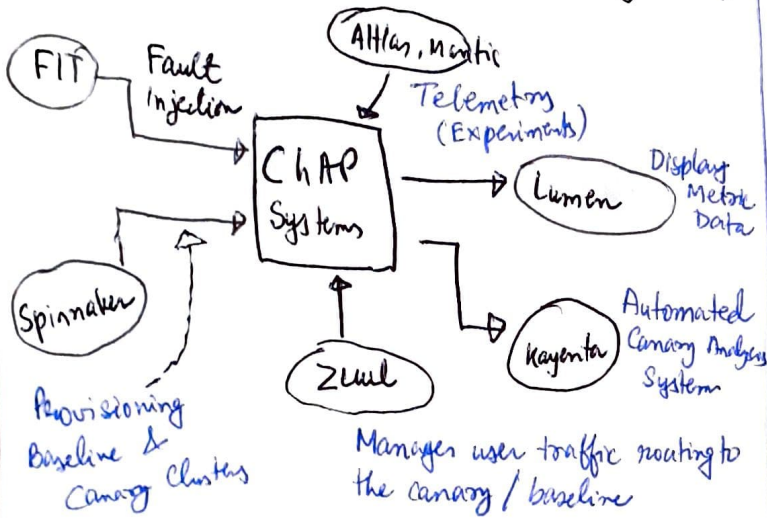
Key Takeaways : Rethinking Approaches

- Use write-ups as learning opportunity.
- Unify all phases of chaos together w/ incidents as well.

ChAP: Automating chaos experiments in Production

Reference: Bastin et al. (ICSE 2019)

- ChAP is developed on the public cloud as a complex set of interacting microservices.
- Two focused failure modes
 - o Services becoming slower (increase in response latency)
 - o Services failing outright (returning errors)



Safeguards

(To minimize the blast radius of the experiment)

- Experiments are run only on weekdays during business hours.
- Not run in failed-over scenarios for independent reasons.
- Set of concurrently running ChAP experiments shouldn't impact over $x\%$ of the N/w traffic.
- If the impact is over threshold, the experiments stop immediately.

Error Rates

Four golden signals (SRE) →

↳ Latency, Throughput, Error Rate, Saturation

Take Aways

- ChAP is excellent as a fault-injection platform
- ChAP for Canary development & load testing experiments.

(Statistical Significance) \neq Validity

- o Accepting A/B test with (say, 95% stat. significance) doesn't correlate with guaranteed uplift. So, stat. significance may not correlate to validity.
- o Stat. significance alone isn't sufficient in deciding when to terminate experiment: Sufficiently long running experiments (even with huge data volume) is an interesting indicator. Important decision parameters →
 - o Includes 1-2 business cycles
 - o Have enough absolute conversions/transactions
 - o Have enough confidence rate (uplift)

Golden Stopping Rules

- Sufficient test durations (1-2 business cycles)
- Well-defined minimum pre-calculated sample size
- Stat. Significance ($\leq 95-99\%$)
- Abs. Number per variation.

- ⊗ Search for uplifts within segment which do not have statistical validity
↳ Destined to Fail

The Art & Science of Interview Engineering

- o Optimality v/s Speed
- o Hack edge cases scenarios (Need to work on that)

Panel : Microservices - Are they still worth it?

What I wish I knew

- ① How to control business complexities with scalable microservices?
- ② One can solve the tradeoff w/ Monolithic architectures (Computational Overhead, No Fault Tolerances) without shifting to Microservices.
- ③ Resolve issues around the architectures (dependencies, simplifying complexities) design w/ Microservices?
- ④ Importance of dev tooling w/ Microservices.

Are they still worth it?

- Y ① Isolation of services, Delivery speedups.
- Y ② Work-independence. Flexible innovation w/o introducing deep complexities.
- N ③ Operational overhead w/ monolithic services offer scalable, cheaper, implementable solution.
- Y ④ Decoupled environment, release processes fits elegantly into the Agile Ecosystems.
 - o Faster change at scalable separation.

What journey Enterprises need to go through for adopting Microservices (Operational Challenge)

- A Technology & Design decision alongside team structures.
- B Monolithic Architecture suffers from low scalable in-memory solution → REDIS is not perfect. (Caching benefits are much more w/ Microservices over Monolithic Architectures)

③ Strong platform team providing efficient platform & toolboxes may become a pre-requisite for transition to Microservices

④ Standards & Conventions are much simpler with Monolithic architectures.

Audience Questions

① Defining bounded context with 1000s of Microservices

(A) Atomic thinking to delegate responsibilities (with minimal dependencies) to microservices.

② Modular Code v/s Microservices

- (A) Good when simplifying the applications & decoupling services.
- (B) Maintenance & Code-Sanity be at severe risk.

Advice on Adoption/Maturing Microservices

- (A) Managed Archive & knowledge base.
- (B) - (i): Well-defined code ownership responsibility
- (ii): Software are not designed in isolation
- (iii): Active communication get Agile.
- (C) Have good idea of the design philosophies & engineering decisions made.
- (D) Keep interested in active toolboxes & platform to support the microservices.

Conway's Law

"Organizations which design systems are constrained to produce designs that are copies of the communication structures of these organizations".